



Precedent Results

Objective

Has anyone already built a commercial product, or filed a patent beyond US12561695B2 or the Agentic Commerce Blueprint research reference (github.com/dmsfiris/agentic-commerce-blueprint), that acts as a merchant-side or payment-processor-side firewall: detecting when product/catalog content shown to an AI shopping agent has been altered, injected, or is stale relative to the last verified state, and blocking the payment-provider call before an order is created? I'm specifically checking for a deployed product, not just protocol standards for agent identity or spend-mandate verification (Visa Trusted Agent, Mastercard Agent Pay, AP2, ACP) — this is about catalog/content integrity at the moment of payment, not who the agent is or how much it's allowed to spend. List everything found, dated, with links, and state plainly whether this specific niche is open or occupied.

Results

No. ## Bottom line

I found no publicly documented, commercially deployed product that occupies the exact niche you defined. I also found no additional patent beyond US12561695B2 that clearly claims the complete combination. On the public evidence reviewed, the niche is therefore **commercially open, but with nearby patent prior art and rapidly emerging research.** This is a public-landscape conclusion, not a freedom-to-operate opinion: unpublished applications, confidential pilots, internal controls, and products described only in sales material may not be visible.

An exact match was required to: 1 operate merchant-side or payment-processor-side; 2 compare the product/catalog representation exposed to an AI shopping agent with a previously verified or attested state, or detect its alteration, injection, or staleness; 3 recheck at checkout/payment time; and 4 stop the PSP/payment-provider invocation **before** order creation. Generic AI firewalls, buyer-side agent guards, feed monitoring, transaction-fraud scoring, and identity/spend-mandate protocols do not meet that definition.

Item	Key dates	Type/status	What it does	Exact-match verdict and missing elements	Link
Exact deployed commercial product	Search status through current review	Commercial product	No deployed merchant-side or payment-processor-side product was found in the gathered evidence that verifies AI-shown catalog/product content against a last verified state and blocks the PSP/payment call before order creation	None found. Niche remains unoccupied in the retrieved evidence; no deployment evidence surfaced for an exact match ([1.1], [1.2], [2.1], [1.3])	N/A
Exact additional patent beyond US12561695B2	Search status through current review	Patent	No additional patent was found in the retrieved patent set that matches all required elements	None found. Closest patents miss at least one of: merchant/processor locus, prior verified catalog-state comparison, or pre-order PSP-call blocking ([3.1], [4.1], [4.2], [3.2], [5.1])	N/A

Item	Key dates	Type/status	What it does	Exact-match verdict and missing elements	Link
US7349871B2	Provisional 2002-08-08; filed 2003-07-29; published 2005-02-03; granted 2008-03-25	Granted US patent	Trusted third party / Secure Transaction Server receives independently generated consumer and merchant views of a purchase agreement, verifies identities and that details are consistent, then takes action to execute the purchasing agreement if conditions are satisfied ([3.1], [3.2])	Near match only. Has agreement-view consistency gating, but not AI shopping agents, not catalog/content integrity vs last verified state, and no evidence it specifically blocks a PSP call before order creation ([3.1], [3.2])	https://patents.google.com/patent/US7349871B2
US7801826B2	Provisional 2002-08-08; filed 2003-07-29; published 2004-05-20; granted 2010-09-21	Granted US patent	System for conducting an agreement via a trusted third party; first and second parties independently generate views of the agreement, third party checks identity and consistency, then executes if conditions are satisfied ([5.1])	Near match only. Same core limitation as US7349871B2: not AI-agent catalog integrity, not stale/injected product-state verification, and no explicit merchant/processor pre-order PSP firewall behavior ([5.1])	https://patents.google.com/patent/US7801826B2
US20260067335 A1	Priority 2024-06-25; filed 2025-11-09; published 2026-03-05	Pending US patent application	Security and privacy preserving agentic browser; detects prompt injection in page content, sanitizes context, validates tool calls with capability broker, requires step-up auth/co-approval for risky actions including transactions, and can deny or defer operations such as order submission ([4.1], [4.3], [4.2])	Near match only. Strong on prompt-injection detection and transaction gating, but it is buyer/browser-side, not merchant-side or processor-side; it does not compare catalog content to a prior verified merchant state; and gathered evidence does not show specific blocking of a PSP call before order creation ([4.1], [4.2])	https://patents.google.com/patent/US20260067335A1
Agentic Commerce Blueprint / A Decision-Centered Reference Architecture for Trustworthy Agentic Commerce	arXiv 2026-07	Research reference / reference implementation	Merchant-side decision architecture proposing protected dependencies, execution-time revalidation, and governance to detect stale or false commercial premises; explicitly frames merchant-side decision ownership as unresolved by ACP/AP2/UCP/MCP and related rails ([1.1], [1.2], [1.3])	Closest conceptual match, but not a deployed product. Explicitly research-stage; evaluation does not establish production security, live PSP interoperability, or deployment ([1.1], [1.3])	https://github.com/dmsfiris/agentic-commerce-blueprint ; https://doi.org/10.48550/arxiv.2607.18347
PayBench	Draft / proposal; retrieved as research proposal with preliminary results	Benchmark proposal	Benchmark for whether AI agents with delegated payment authority should attempt payment at all; studies unsafe commercial autonomy, including prompt injection in merchant/tool outputs competing with spending policy ([6.1])	Not a product; not a patent. Proposal only; no merchant-side deployed firewall, no prior-state catalog verification, no PSP interposition product evidence ([6.1])	https://www.arxiv.org/overview/2510.19997v2
RAILS: Verification-Native Clearing for Agentic Commerce	arXiv 2026-06-07	Research protocol / clearing proposal	Clearing layer for agentic commerce; separates authorization, execution, performance, attribution, loss, and settlement; explicitly says payment/authorization rails do not determine whether an obligation was actually satisfied ([7.1], [7.2])	Not an exact match. It is a verification-native clearing proposal, not a merchant/processor catalog-integrity firewall that compares AI-visible catalog state and blocks the payment-provider call before order creation ([7.1], [7.2])	https://doi.org/10.48550/arxiv.2606.08790
PCAT / AIP-Bench (Protocol-Level Attacks on Agentic Commerce Platforms)	arXiv 2026-07	Academic taxonomy / benchmark / unified defense	Comparative academic work on protocol-level vulnerabilities in agentic commerce; states no prior work systematically addressed these structural vulnerabilities and notes production agent clouds still lack mature controls ([2.1])	Not a deployed product. Reference-defense/benchmark framing only; no evidence of commercial merchant/processor deployment for catalog-state verification plus pre-order PSP blocking ([2.1])	https://doi.org/10.48550/arxiv.2607.21824
Excluded category: Visa Trusted Agent, Mastercard Agent Pay, AP2, ACP	Contemporary standards/protocols; dates not material here	Excluded standards/protocols	Identity, mandate, authorization, or agent/payment protocol layers; complementary but not the target niche ([1.1], [1.2], [7.1])	Excluded, not exact candidates. They address who the agent is or what it may spend, not merchant-side catalog/content integrity at payment time relative to a last verified state ([1.1], [1.2], [7.1])	https://doi.org/10.48550/arxiv.2607.18347 ; https://doi.org/10.48550/arxiv.2606.08790

Table: This table summarizes the materially relevant product, patent, and research findings against the exact niche definition. It distinguishes true matches from near matches and excluded standards, with dates, links, and the specific missing elements for each.

Closest patents—but not occupancy of the niche

1. Fujitsu US7349871B2

Priority: August 8, 2002. **Filed:** July 29, 2003. **Published:** February 3, 2005. **Granted:** March 25, 2008. [Google Patents](#).

This patent has a genuinely relevant primitive: a consumer and merchant independently generate respective “views” of a purchasing agreement; a trusted Secure Transaction Server checks identities and whether the views are consistent; it takes action to execute the agreement only if the conditions are satisfied. ([3.1], [3.2])

It nevertheless is **not the claimed niche**. It does not address an AI shopping agent, catalog poisoning or prompt injection, staleness relative to the last verified catalog state, or a merchant/processor interceptor explicitly blocking the PSP call before order creation. It is best treated as broad prior art concerning independent transaction-view reconciliation, not as an agentic catalog-integrity firewall.

2. Fujitsu US7801826B2

Priority: August 8, 2002. **Filed:** July 29, 2003. **Published:** May 20, 2004. **Granted:** September 21, 2010. [Google Patents](#).

This related system patent likewise has two parties independently generate views of an agreement; the trusted third party checks their consistency and executes only when conditions are satisfied. ([5.1])

The same exclusions apply: no AI-facing catalog baseline, no alteration/injection/staleness comparison, and no explicit pre-order PSP-call firewall. It is relevant prior art for “two-view consistency before transaction execution,” but it does not occupy the narrower implementation described in the question.

3. US20260067335A1, “Security and Privacy Preserving Agentic Browser”

Priority: June 25, 2024. **Filed:** November 9, 2025. **Published:** March 5, 2026. **Status:** pending. [Google Patents](#).

This is the closest agent-specific patent found. Its claims cover separating trusted instructions from untrusted web content, detecting prompt-injection indicators, sanitizing context, validating tool calls through scoped capability tokens, and gating sensitive calls that initiate transactions. It also claims simulation and blocking upon a policy violation. ([4.1])

The specification gives an Amazon-order example: inspect page content for prompt injection, require authentication and user approval, enforce spend/scope limits, and withhold order submission if checks fail. It also describes blocking form submission when a coupon site embeds instructions seeking saved card data. ([4.3], [4.2])

It still **does not match** because it is principally **buyer/browser-side**, detects hostile instructions rather than comparing catalog fields with a merchant’s last verified state, and does not specifically place a merchant/processor control immediately before the PSP invocation and before order creation. No commercial deployment evidence was found.

Research systems and proposals

The [Agentic Commerce Blueprint](#) remains the closest direct conceptual reference. Its associated paper describes protected dependencies and execution-time revalidation for stale or false commercial premises, while expressly framing merchant-side commercial-meaning enforcement as unresolved by ACP, AP2, UCP, MCP, wallets, and payment networks. Its evaluation does **not** establish production security, live payment-provider interoperability, or commercial deployment. ([1.1], [1.2], [1.3])

[PayBench](#) asks the adjacent question of whether an agent should attempt a payment when merchant content or tool output contains prompt injection. It expressly describes itself as a research proposal with only preliminary results and a planned mock-merchant/payment environment. It is neither a deployed firewall nor prior-state catalog verification. ([6.1])

RAILS, dated June 7, 2026, proposes verification-native clearing for agentic commerce. It correctly distinguishes payment authorization from determining whether an obligation was fulfilled and can produce settlement instructions from evidence. However, its object is generalized obligation clearing/performance verification, not comparison of AI-visible catalog content to a last verified merchant state immediately before a PSP call. It is a protocol paper, not evidence of a deployed commercial product. ([7.1], [7.2])

PCAT/AIP-Bench, dated July 2026, is an academic protocol-attack taxonomy, benchmark, and reference defense. It reports that structural protocol vulnerabilities had not previously been addressed systematically and that production agent clouds still lacked mature controls. It does not document commercial deployment of the requested merchant/processor firewall. ([2.1])

Why the named payment standards do not change the answer

Visa Trusted Agent, Mastercard Agent Pay, AP2, and ACP concern agent recognition, authorization, mandate binding, credentials, or interoperable transaction messages. The reviewed architecture literature explicitly separates these functions from the merchant's unresolved responsibility to ensure that the current commercial premise is true, fresh, and authorized. ([1.1], [1.2])

Likewise, RAILS describes Visa, Mastercard, AP2, and other rails as answering pieces of authorization or value transfer, while not themselves determining whether the underlying commercial obligation was correctly represented or fulfilled. ([7.1])

Plain conclusion

Commercial-product status: open. No exact deployed product was found.

Patent status beyond US12561695B2: no exact additional filing found. The most relevant nearby patents are US7349871B2/US7801826B2 for reconciling independently generated transaction views and US20260067335A1 for buyer-side prompt-injection detection and transaction gating. None combines the required merchant/processor locus, last-verified catalog-state comparison, checkout-time revalidation, and pre-order PSP-call block.

Accordingly, the narrow product category appears **unoccupied in the public record searched**, although it is not a blank-slate intellectual-property field. The older two-view transaction patents and newer agentic-browser security filing should be considered in claim drafting and freedom-to-operate analysis, while the Blueprint and later clearing/security papers may be relevant non-patent literature.

References

[1.1] A Decision-Centered Reference Architecture for Trustworthy Agentic Commerce. Dimitrios S. Sfiris. ArXiv (2607). <https://doi.org/10.48550/arxiv.2607.18347> Context: "This excerpt from a design-science paper on agentic commerce presents a reference architecture for the merchant-side decision layer. The author explicitly identifies that existing protocols (ACP, AP2, UCP, MCP) address complementary integration aspects but are not the merchant's complete internal decision model. Critically, the paper states that current findings "leave a separate merchant-side question unresolved: how should a commerce system represent and enforce the commercial meaning behind an agent-facing response?" The paper proposes the Agentic Commerce Blueprint as a solution for detecting state changes (price, inventory, policy, checkout, mandate, or evidence changes) and preventing use of stale decisions through protected dependencies, execution-time revalidation, and generated-claim governance. However, the abstract explicitly states the evaluation "do not establish rule completeness, production security, performance, legal compliance, live interoperability, population error rates, or independent replication," indicating this is research-stage work, not a commercially deployed product. No evidence of deployed merchant-side or payment-processor-side products is presented."

[1.2] A Decision-Centered Reference Architecture for Trustworthy Agentic Commerce. Dimitrios S. Sfiris. ArXiv (2607). <https://doi.org/10.48550/arxiv.2607.18347> Context: "The excerpt presents a reference architecture for trustworthy agentic commerce but explicitly states this is a reference implementation and semantic contract, not a production-deployed system. The paper identifies unresolved gaps in merchant-side decision ownership across existing protocols (ACP, AP2, UCP, MCP, wallets, payment networks). It emphasizes that protocol messages can be syntactically valid while their commercial premise is false, stale, or unauthorized—the core problem being semantic drift across boundaries. The work acknowledges that agentic payment security cannot rely solely on signing mandates and that authority artifacts must be evaluated against current commercial state rather than treated as final permission. The excerpt indicates this is foundational research addressing protocol limitations rather than reporting an existing deployed solution for detecting altered or stale catalog content before payment-provider invocation."

[1.3] A Decision-Centered Reference Architecture for Trustworthy Agentic Commerce. Dimitrios S. Sfiris. ArXiv (2607). <https://doi.org/10.48550/arxiv.2607.18347> Context: "This excerpt describes a reference architecture (Agentic Commerce Blueprint v0.9.2) designed to address merchant-side decision integrity during checkout, including detection of altered or stale catalog content through dependency revalidation and state integrity controls. However, the authors explicitly exclude vendor announcements lacking inspectable technical content and note their evaluation does not measure live protocol and payment-provider interoperability. The artifact is a reference implementation and semantic evaluation, not a deployed commercial product. The authors present this as research addressing unresolved gaps in merchant-side decision ownership—specifically replay, context binding, protocol conformance, and state drift threats. No commercially deployed merchant-side or payment-processor product is identified in this excerpt."

[2.1] Protocol-Level Attacks on Agentic Commerce Platforms: A Cross-Platform Taxonomy, AIP-Bench, and Unified Defense. Yedidel Louck. ArXiv (2607). <https://doi.org/10.48550/arxiv.2607.21824> Context: "The excerpt does not provide evidence of a commercially deployed merchant-side or payment-processor-side product that detects altered, injected, or stale catalog/product content and blocks payment calls. The paper explicitly acknowledges that "no prior work" has systematically identified and addressed structural protocol-level vulnerabilities in agentic commerce platforms through a multi-platform comparative study. The excerpt indicates that "production agent clouds still lack mature registry and discovery controls," suggesting the problem remains unresolved in deployed systems. The paper discusses verification-native clearing and execution-bound primitives as research proposals to bind payment authorization to execution, but these are presented as literature proposals rather than deployed commercial solutions. PCAT is introduced as a unified defense mechanism, but it is framed as an academic reference implementation addressing protocol-layer concerns rather than a merchant or payment-processor product."

[3.1] Methods for purchasing of goods and services. Yannis Labrou. Patent (US) (2008). <https://patents.google.com/patent/US7349871B2/en> Context: "US7349871B2 discloses a method for conducting purchasing agreements between consumers and merchants through a trusted third party using wireless networks. The method involves: (1) the consumer generating and transmitting a first view of the agreement to the third party, (2) the merchant independently generating and transmitting a second view of the agreement to the third party, (3) the third party receiving both views and verifying the identities of both parties, (4) confirming that the details of the independently generated views are consistent, and (5) executing the purchasing agreement if conditions are satisfied. The third party includes a Secure Transaction Server. The patent was filed July 29, 2003, with a provisional application filed August 8, 2002, and was granted March 25, 2008. The excerpt provided does not contain specific claim language beyond the abstract description, nor does it discuss payment services contact procedures."

[3.2] Methods for purchasing of goods and services. Yannis Labrou. Patent (US) (2008). <https://patents.google.com/patent/US7349871B2/en> Context: "# Summary This excerpt presents US Patent 7,349,871 B2 (filed July 29, 2003; issued March 25, 2008), which describes "Methods for Purchasing of Goods and Services." The patent, assigned to Fujitsu Limited, covers a purchasing agreement system involving a consumer, merchant, and "trusted third party" using a wireless network. The method includes: generating independent views of an agreement by consumer and merchant, transmitting both to a third party, verifying identities, and confirming "that the details of the independently generated views of the agreements are consistent" before executing the purchase. The system utilizes a "Secure Transaction Server" as the central verification component. However, this patent focuses on agreement consistency verification through a trusted intermediary, not specifically on detecting altered/injected product catalog content or preventing payments based on catalog integrity validation at the moment of transaction."

[4.1] Security and Privacy Preserving Agentic Browser. Bao Tran. Patent (US) (2026). <https://patents.google.com/patent/US20260067335A1/en> Context: "# Summary This excerpt describes a patent application (US20260067335A1, filed 2025-11-09, published 2026-03-05) for an "Security and Privacy Preserving Agentic Browser" with inventors including Bao Tran and assignees including Stripe, Inc., WithSecure Corporation, and Dell Products L.P. The patent addresses AI browser risk governance through: (1) classifying proposed actions as safe/risky using LLM weights or policy rules; (2) initiating step-up authentication for risky actions; (3) preventing prompt injection by separating trusted instructions from untrusted content and sanitizing via redaction/summarization; (4) validating tool calls against capability tokens scoped to domain, intent, and time windows; (5) enforcing human-in-the-loop gates for credential access, transactions, or data exfiltration; (6) masking sensitive surfaces (credential forms, payment instruments, account numbers); and (7) maintaining device-local encrypted vector stores excluding sensitive-domain embeddings from remote inference. The patent does not explicitly address catalog/content integrity verification or detecting altered product information at payment moments."

[4.2] Security and Privacy Preserving Agentic Browser. Bao Tran. Patent (US) (2026). <https://patents.google.com/patent/US20260067335A1/en> Context: "# Summary The excerpt describes a comprehensive security framework for AI agents performing sensitive actions, including an Amazon order placement example. Key features include: **Detection mechanisms:** An "injection detector" inspects page content for prompt injection vectors in DOM, CSS, and image alt text, producing "sanitized context" by redacting suspicious tokens. **Transaction controls:** Step-up authentication (passkey/biometric) is required before credential use or payment submission. A "capability broker" mints timeboxed tokens bound to origin, session, and specific action. **Action verification:** System presents user-facing summaries listing items, destination, cart total, required capabilities, and risk indicators before requiring co-approval within a "bounded time window." **Enforcement:** Spend/scope limits are checked pre-execution; clipboard and cross-origin access are blocked during sensitive surfaces. Execution uses "write-only autofill" so credentials are never read back into model context. **Logging:** An "immutable, append-only audit log" records transaction details, risk scores, approvals, and outcomes. The

framework emphasizes preventing injection-based manipulation and ensuring user approval before payment execution, but does not specifically address catalog integrity verification or detecting altered product content at payment time."

[4.3] Security and Privacy Preserving Agentic Browser. Bao Tran. Patent (US) (2026). <https://patents.google.com/patent/US20260067335A1/en> Context: "# Summary The excerpt describes a system for securing AI agent interactions with web pages, particularly for sensitive actions like online shopping. Key mechanisms include: **Content Analysis (S500)**: Detects sensitive surfaces (credentials, financial identifiers, personal data) using DOM analysis, rule engines, and classifiers, assigning confidence levels and sensitivity classes while suppressing non-sensitive elements. **Embeddings & Local Storage (S600-S602)**: Creates device-local embeddings from sanitized browsing context with sensitive-surface masking, credential exclusion, and risk scoring. Vectors are encrypted, domain-partitioned, and stored locally with content hashes and provenance indicators. **Access Control (S604)**: Retrieves vectors restricted to domain partitions within trusted execution environments using per-domain keys; cross-domain retrieval is disabled by default. **Governance (S700-S708)**: Classifies elevated-risk actions, initiates step-up authentication, presents action summaries for co-approval, enforces user-configured spend/scope limits, and records immutable audit logs with timestamps and masked sensitive values. **Shopping Workflows (S304, S506)**: Separates read-only steps from fund-affecting actions; requires user interlocks and spend-limit confirmation before submission for purchasing agents. The system emphasizes local processing, encryption, and human approval gates without exposing raw confidential material."

[5.1] Framework and system for purchasing of goods and services. Yannis Labrou. Patent (US) (2010). <https://patents.google.com/patent/US7801826B2/en> Context: "# Summary This excerpt presents US Patent 7,801,826 B2 (filed July 29, 2003; issued September 21, 2010), which describes "a system for conducting an agreement between two parties relying on a trusted third party." The patent details a framework where parties independently generate views of agreements, transmit them to a trusted third party through wireless/wired networks, and the third party "verifies conditions including that the identities of the parties that transmitted the agreements and that the independent views of the agreement are consistent with each other, and takes action to execute the agreement if the conditions are satisfied." The patent covers general transaction frameworks with merchant and financial institution servers but does not specifically address catalog/content integrity verification, AI shopping agents, payment processor firewalls, or detection of altered product data at the moment of payment. It focuses on agreement verification between parties rather than the specific niche described in your question."

[6.1] PayBench: A Benchmark for Unsafe Commercial Autonomy in AI Agents with Delegated Payment Authority. C Plunkett. <https://paybench.org/paybench.pdf> Context: "# Summary The excerpt presents PayBench, a research benchmark for measuring AI agent safety under delegated payment authority. It focuses on whether agents preserve user intent when executing payments, not authorization infrastructure. PayBench addresses a gap: existing benchmarks (τ -bench, ToolEmu, AgentDojo, InjecAgent) measure agent capability or tool security, but "none isolate the decision to pay under delegated authority." The benchmark uses "matched safe-to-act and unsafe-to-act scenario pairs" to prevent reward-hacking through blanket refusal, and identifies "twelve failure modes grouped under five control categories (spend limits, authorization scope, consent and escalation, privacy and disclosure, and adversarial robustness)." The research is explicitly described as a "proposal" with "only preliminary results exist"—it is not deployed. No commercial products, patents beyond those mentioned in your question, or deployed systems for detecting catalog/content integrity at payment time are identified in this excerpt."

[7.1] RAILS: Verification-Native Clearing For Agentic Commerce. Adrian de Valois-Franklin, Alexandrescu Bogdan. ArXiv (2606). <https://doi.org/10.48550/arxiv.2606.08790> Context: "# Summary This excerpt introduces RAILS (Real-Time Agent Integrity & Ledger Settlement), a framework addressing the "agentic clearing problem" in autonomous agent commerce. The authors argue that existing infrastructure—tool protocols (MCP), agent communication (A2A), payment rails (x402), and mandate frameworks (AP2, Visa Trusted Agent, Mastercard Agent Pay)—are insufficient because none determines whether delegated obligations were fulfilled. RAILS proposes seven primitives enabling verification graded by evidence admissibility, with a "soundness property: no financially material settlement is supported by evidence below the obligation's admissibility floor."

The paper distinguishes clearing from authorization, payment, escrow, and LLM-as-judge evaluation. It addresses six separable transaction questions: authorization, execution, performance, attribution, loss, and settlement. The authors state they are "not aware of a prior agent-commerce verification mechanism that states a property of this kind."

[7.2] RAILS: Verification-Native Clearing For Agentic Commerce. Adrian de Valois-Franklin, Alexandrescu Bogdan. ArXiv (2606). <https://doi.org/10.48550/arxiv.2606.08790> Context: "# Summary The excerpt describes RAILS, a protocol framework for verifying AI agent outputs through multiple verifier classes. Key relevant points: **Verifier Classes (Section 5.6):** Five reference verifier types assess outputs: constraint (deterministic predicates), receipt (external system receipts), semantic (LLM-based evaluation), policy (compliance), and human (arbitration). However, the excerpt identifies a "structural blind spot"—none assess "the behavioral reliability of the model that produced it." **Proposed Solution:** A new verifier class combining "Ranking Inference verification primitive" and "Machine Psychometrics behavioral profiling" to measure model reliability, sycophancy, and perturbation instability. This becomes the "RAILS Score," an always-on triage verifier. **Lifecycle (Section 6):** Ten-stage process from Discovery through Settlement Instruction to Finality, with verification, adjudication, and appeal stages. **Threat Model (Section 5.7):** Three attack families: FORGE-UP (claim false evidence class), LAUNDER-BASIS (misreport verifier basis), DOWNGRADE-FLOOR (weaken obligation binding). The excerpt does not describe a deployed commercial product detecting altered catalog content before payment processing."